

Moderator-Assisted Granger Causal Network Learning for High-Dimensional Multi-Subject Longitudinal Data

José Á. Sánchez Gómez Holly O’Rourke Gene Brewer *

Abstract

In neuroscience and behavioral psychology, researchers are interested in using longitudinal data to estimate the causal effects that variables have on each other over time. These relationships can be modeled via Granger-causal networks (GCN), where nodes correspond to variables, and directed links indicate lagged temporal relationships among them. When longitudinal data is gathered over multiple subjects, GCN models must account for potential subject heterogeneity. While methods in the literature can perform estimation of heterogeneous GCNs, they often lack interpretable explanations of the source of such measured heterogeneity. In the present paper, we propose the moderator vector auto-regressive model (MOD-VAR) for interpretable estimation of heterogeneous GCN in multi-subject longitudinal data. We achieve this by using external moderator variables to explain the heterogeneity of GCNs across subjects. Our theoretical results confirm the improved estimation of MOD-VAR for heterogeneous GCN estimation in high dimensional scenarios, even when the GCN are highly heterogeneous. Furthermore, our extensive numerical simulations confirm the superior performance of our proposed MOD-VAR for GCN estimation and forecasting in a wide variety of settings. *We analyze the XXXX dataset, and find that heterogeneity in behavioral dynamics can be partially explained by XXX and YYY.*

Keywords: Adaptive estimation, network regression, time series, vector auto-regressive model, weighted networks.

*José Á. Sánchez Gómez is Assistant Professor, Department of Statistics, University of California Riverside, Riverside, CA 92521, USA. E-mail: josesa@ucr.edu. Holly O’Rourke is Associate Professor, Department of Psychology, University of California Riverside, Riverside, CA 92521, USA. Email: holly.orourke@ucr.edu. Gene Brewer is Professor, Department of Psychology, University of California Riverside, Riverside, CA 92521, USA.

Corresponding Authors: José Á. Sánchez Gómez (josesa@ucr.edu), Holly O’Rourke (holly.orourke@ucr.edu).

1 Introduction

The modeling and forecasting of intensive multivariate time series data has been a long-standing area of research in statistics, optimization and machine learning [9], with applications in economics, finance, neuroscience and psychology [20; 11; 21; 4; 12], among other areas. In particular, researchers are often interested in constructing a *Granger-causal network* (GCN) to describe the dependence that variables in the time series have over time. In a GCN, nodes correspond to variables in the time series, and a directed link between two nodes represents an effect of the origin variable in the present to future values of the target variable. When the time dependence is linear, the GCN can be fully recovered by estimating coefficients of a sparse vector auto-regressive model (VAR). This has led to the development of a wide variety of methods aimed at accurately estimating GCNs in time series data [1; 2; 19]. For a review, see Wei [26].

One major statistical challenge is that of estimating GCNs over multiple subjects. For example, in fMRI studies, researchers are often interested in determining the patterns of connectivity among different regions of the brain [11; 21; 4]. By processing functional magnetic resonance imaging (fMRI) data, researchers can measure the activity of brain regions over time. In this context, a GCN may describe the connectivity of electrical signals across different brain regions. In multi-subject studies, we expect the connectivity patterns to have shared components, while also showing some heterogeneity across different subjects. In this setting, fitting a single VAR model may lead to bias and inaccurate forecasting. On the other hand, training separate VAR models for each subject may miss the chance to exploit the common structure to obtain improved forecasting accuracy [5].

Several works have explored the problem of fitting and testing models for multi-subject high dimensional GCNs. The Multi-VAR method [6; 7] applies penalized estimation to recover both joint and individual components of the GCNs across multiple subjects. The group iterative multiple model estimation procedure (GIMME) decomposes shared and individual components of both a GCN and structural equation model across subjects [8; 12; 13]. Recent works have explored the task of testing the significance of sparse GCN estimators. For example, Zhu and Liu [29]; Uematsu and Yamagata [23] develop multiple testing procedure based on the debiased lasso [24; 27; 10] to determine if there are significant

differences in the GCN across multiple subjects.

Despite the progress of prior works, a key question remains unanswered: *can we explain the heterogeneity observed across subjects?* For example, in the aforementioned example of fMRI data, researchers may be interested not only in detecting heterogeneity in the GCNs across individuals, but to determine whether external moderator like age, gender, and existing health conditions or treatments may explain it. While existing methods measure GCN heterogeneity, they are ill-equipped to determine if the variation in the GCNs can be explained by such external moderator data.

While models that link the GCN of time series with moderator variables have been explored via multi-level modeling [22; 14], such methods have limited applicability since they are unable to handle a large number of variables d or moderator variables p unless the number of subjects or time-measurements per subjects are prohibitively large. Alternatively, the VARX model incorporates the effect of external moderators to the VAR forecasting model [19; 18]. Despite its advantages, the external moderators are only considered as potentially affecting the mean behavior of the time series. From this, the VARX does not consider the potential effect that moderators have on the GCN. Thus, there is a pressing need for methods that can estimate the influence that moderator variables have on the GNC across multiple subjects, and can provide reliable estimates, even in cases where the number of parameters may be significantly larger than the number of observations.

In the present work, we propose the moderator vector autoregressive model (MOD-VAR) for the simultaneous estimation of GCN over multi-subject longitudinal data. Our method measures the effect that external moderators may have on the GCNs across datasets. We achieve this by decomposing each GCN into (i) connections shared across all datasets, (ii) connections in the network dependent on external moderators, and (iii) connections that are fully individual, and cannot be explained by the given moderator variables. Our method is highly general, encompassing the VAR and Multi-VAR framework [6; 7]. While similar moderator models have been explored in the context of Gaussian graphical models [17; 28; 15], to our best knowledge this is the first method to exploit the presence of shared, moderated and individual auto-regressive structure for enhancing the forecasting and interpretability of multi-subject high-dimensional GCNs.

Our theoretical findings confirm that our proposed MOD-VAR can achieve consistent

estimation of all parameters, as long as our total sample size N satisfies $N \gg \log(pd^2)$. Furthermore, we conduct extensive simulation experiments that further confirm the improved performance of our proposed MOD-VAR when compared to other methods designed for the estimation of common and individual components of GCNs for multiple subject time series data. We obtain improved performance in scenarios in which the heterogeneity is fully explained by the moderators, partially explained by the moderators, and even when heterogeneity is unrelated to the moderators. To showcase the usefulness of our method in real data applications, we analyze a dataset of brain activation across multiple regions via fMRI measurements, and quantify the effect that external moderators like age, gender and size of the locus coeruleus (LC) region has on brain connectivity across subjects.

The rest of the paper is organized as follows. In section 2 we introduce our model of moderator time series model, and introduce our MOD-VAR estimator. In section 3 we provide details on the implementation of our proposed method. In Section 6 we provide numerical comparisons of our proposed MOD-VAR with methods found in the literature. In Section 7 we exemplify the use of our MOD-VAR method in the analysis of daily mood and behavior psychological data. Closing remarks and discussions are found in Section 8.

Notational Conventions

We establish the following notational conventions. For any n , let $[n] = \{1, 2, \dots, n\}$. For any $k \leq \ell$, we denote $[k : \ell] = \{k, k + 1, \dots, \ell - 1, \ell\}$. For any $A \in \mathbb{R}^{p \times q}$, we denote the transpose of A as A' . For $\mathbf{x} = (x_1, x_2, \dots, x_d)' \in \mathbb{R}^d$, we denote by $\|\mathbf{x}\|_1, \|\mathbf{x}\|_2, \|\mathbf{x}\|_\infty$ the ℓ_1 -norm, euclidean norm and entrywise maximum norm respectively, given by $\|\mathbf{x}\|_1 = \sum_{i=1}^d |x_i|$, $\|\mathbf{x}\|_2 = \sqrt{\mathbf{x}'\mathbf{x}}$, and $\|\mathbf{x}\|_\infty = \max_{i \in [d]} |x_i|$. For any $\Sigma = (\sigma_{ij})_{i,j=1}^d \in \mathbb{R}^{d \times d}$, we denote the trace as $\text{tr}(\Sigma) = \sum_{i \in [d]} \sigma_{ii}$. For $A \in \mathbb{R}^{p \times q}$, we denote by $\|A\|_2$, and $\|A\|_F$ the ℓ_2 operator, and Frobenius norms, respectively, given by $\|A\|_2 := \sup_{\mathbf{x} \in \mathbb{R}^q \setminus \{\mathbf{0}\}} \|A\mathbf{x}\|_2 / \|\mathbf{x}\|_2$, and $\|A\|_F = \sqrt{\text{tr}(A'A)}$. Additionally, we denote by $\|A\|_1$ the vectorized ℓ_1 -norm of the matrix A , given by $\|A\|_1 = \sum_{i,j} |A_{ij}|$. For any k , we denote $\mathbf{1}_k \in \mathbb{R}^k$ the vector with all one entries. For any set A , we denote by $\mathbf{1}(A)$ the indicator function of A . For any $x \in \mathbb{R}$, we denote $\text{sign}(x) = \mathbf{1}(x > 0) - \mathbf{1}(x < 0)$. For any symmetric matrix $\Sigma \in \mathbb{R}^{d \times d}$, we denote the minimum and maximum eigenvalues of Σ as $\Lambda_{\min}(\Sigma), \Lambda_{\max}(\Sigma) \in \mathbb{R}$, respectively.

2 Moderator Vector Autoregressive Model

2.1 Vector Autoregressive Models

Consider a study with n subjects. For each subject $k \in [n]$, we have access to the measurements of d different variables of interest over T_k equally spaced time points, which we denote by $\mathbf{Y}_k = \{\mathbf{Y}_t^{(k)}\}_{t=1}^{T_k} \subset \mathbb{R}^d$. We model the lagged temporal relationships that exist between the different variables, by a VAR(1) model, expressed as

$$\mathbf{Y}_t^{(k)} = \Phi^{(k)*} \mathbf{Y}_{t-1}^{(k)} + \boldsymbol{\varepsilon}_t^{(k)}, \quad \text{for all } 2 \leq t \leq T_k,$$

where $\Phi^{(k)*}$ is a $d \times d$ matrix that represents the true lag-1 multivariate relationships for the subject k . In particular, for $1 \leq i, j \leq d$, the entry $\Phi_{ij}^{(k)*}$ represents the effect that the variable j at time $(t-1)$ affects the variable i at time t . Notice that we define a distinct coefficient matrix $\Phi^{(k)*}$ for each subject $k \in [n]$. Given the matrix $\Phi^{(k)*}$, we define the GCN associated with subject k , as the network with set of nodes $N = \{1, 2, \dots, d\}$, and a link $i \rightarrow j$ is present in the GCN if and only if $\Phi_{ji}^{(k)*} \neq 0$. Thus, the structure of the GCN depends on the sparsity pattern of the matrix $\Phi^{(k)*}$.

For each subject $k = 1, 2, \dots, n$, in addition to the time series data $\mathbf{Y}_k$, we have a set of p potential moderator variables $\mathbf{X}_k = (X_{k1}, X_{k2}, \dots, X_{kp})' \in \mathbb{R}^p$. Notice that the moderators in $\mathbf{X}_k$ are not time-varying, and thus can be thought as stable features of each subject. For example, in a neuroimaging study analyzing brain connectivity, these moderators may include clinical or demographic variables such as gender or age. We may also consider controlled experimental variables of interest such as treatment status of individual, or magnitude of the dose provided to a patient. Our goal in this work is to determine if such moderator variables explain the heterogeneity in the GCNs across subjects, and exploit this potential effects to improve the estimation of the GCNs.

2.2 Moderator Vector Autoregressive Models

We first introduce the scenario in which all cross-subject heterogeneity can be explained by the moderators in $\mathbf{X}_k$. In this case, we can represent the moderator-dependent relationships as $\Phi^{(k)*} = \Phi^*(\mathbf{X}_k)$, where, for any vector of p moderators $\mathbf{x} = (x_1, x_2, \dots, x_p)'$,

$$\Phi^*(\mathbf{x}) := \Phi^{(0)*} + x_1 \Phi^{(M1)*} + x_2 \Phi^{(M1)*} + \dots + x_p \Phi^{(Mp)*}, \quad (1)$$

where $\Phi^{(0)*}, \Phi^{(M1)*}, \Phi^{(M2)*}, \dots, \Phi^{(Mp)*} \in \mathbb{R}^{d \times d}$. The matrix $\Phi^{(0)*} \in \mathbb{R}^{d \times d}$ represents the shared connectivity in the GCN across subjects. For each moderator variable $1 \leq \ell \leq p$, the matrix $\Phi^{(M\ell)*}$ represents the effect that the ℓ -th moderator has on the time series dynamic. More specifically, for $i, j \in [d]$ and $\ell \in [p]$ the entry $\Phi_{ij}^{(M\ell)*}$ represents the effect that the moderator variable ℓ has on the lagged relationship of j on the variable i . If for a subject i we have $x_{i\ell} = 0$, then $x_{i\ell}\Phi^{(M\ell)*} = \mathbf{0}$ and thus, the moderator ℓ does not have any effect on the time dynamics for subject i . In addition, the large values of $x_{i\ell}$ are associated with stronger effects of the moderator ℓ on the dynamic of subject i . Under the model (1), we define the moderator vector auto-regressive (MOD-VAR) model via the equation

$$\mathbf{Y}_t^{(k)} = \Phi^*(\mathbf{X}_k)\mathbf{Y}_{t-1}^{(k)} + \boldsymbol{\varepsilon}_t^{(k)}, \quad \text{for all } 1 \leq k \leq n \text{ and } 2 \leq t \leq T_k. \quad (2)$$

We illustrate our MOD-VAR model from (4) with the following example.

Example 1 (Binary moderators). *We consider a VAR(1) model with n subjects. For each subject, we measure a single moderator variable (i.e., $p = 1$) $X_1, X_2, \dots, X_n \in \mathbb{R}$ with $\{X_k\}_k \stackrel{i.i.d.}{\sim} \text{Bernoulli}(p)$. In this case, our propose model can be written as,*

$$\mathbf{Y}_t^{(k)} = [\Phi^{(0)*} + X_{k1}\Phi^{(M1)*}] \cdot \mathbf{Y}_{t-1}^{(k)} + \mathbf{e}_t^{(k)}$$

Notice that the GCN of a subject $k \in [n]$ incorporates the matrix $\Phi^{(0)} \in \mathbb{R}^{d \times d}$ shared across all subjects. Furthermore, there is an individual structure that depends on the value of X_k . All subjects with $X_k = 1$ will have GCN $\Phi^{(k)*} = \Phi^{(0)*} + \Phi^{(M1)*}$, while subjects with $X_k = 0$ have GCN $\Phi^{(k)*} = \Phi^{(0)*}$. Thus, the matrix $\Phi^{(M1)*}$ represents the structural differences in the GCN across the the grouping given by the moderators $X_1, X_2, \dots, X_n$.*

Now, consider modeling two binary moderator variables $\mathbf{X}_k = (\mathbf{X}_{k1}, \mathbf{X}_{k2})^\top \in \{0, 1\}^2$, we have now a total of four potential GCNs: (i) $\Phi^{(k)} = \Phi^{(0)*}$ if $\mathbf{X}_k = (0, 0)^\top$, (ii) $\Phi^{(k)*} = \Phi^{(0)*} + \Phi^{(M1)*}$ if $\mathbf{X}_k = (1, 0)^\top$, (iii) $\Phi^{(k)*} = \Phi^{(0)*} + \Phi^{(M2)*}$ if $\mathbf{X}_k = (0, 1)^\top$, and (iv) $\Phi^{(k)*} = \Phi^{(0)*} + \Phi^{(M1)*} + \Phi^{(M2)*}$ if $\mathbf{X}_k = (1, 1)^\top$. We illustrate the GCNs of two subjects with different moderator values in Figure 1.*

2.3 MOD-VAR with Fully Individual Connections

While external moderators may account for some of the GCN heterogeneity across subjects, often some heterogeneity may remain unexplained. Thus, in addition to a shared component $\Phi^{(0)*}$, and moderator effects $\Phi^{(M1)*}, \dots, \Phi^{(Mp)*}$, we consider $\boldsymbol{\Delta}^{(I1)*}, \boldsymbol{\Delta}^{(I2)*}, \dots, \boldsymbol{\Delta}^{(In)*} \in$




Figure 1: Illustration of the MOD-VAR framework with $p = 2$ binary moderator variables, as discussed in Example 1. Top-panels: visualization of the decomposition of the GCN of subject 1. Since $\mathbf{X}_1 = (1, 0)^\top$, only the first moderator has an effect on the GCN for this subject. Bottom panels: visualization of the GCN for subject 2. Since $\mathbf{X}_2 = (0, 1)^\top$, only the second moderator affects the GCN structure.

$\mathbb{R}^{d \times d}$ matrices corresponding to fully individual GCN connections. For each subject k , the coefficient matrix $\Phi^{(k)*}$ can be decomposed as

$$\Phi^{(k)*} = \Phi^*(\mathbf{X}_k) + \Delta^{(Ik)*} = \Phi^{(0)*} + \left(\sum_{\ell=1}^p X_{k\ell} \Phi^{(M\ell)*} \right) + \Delta^{(Ik)*}. \quad (3)$$

Notice that $\Delta^{(Ik)*} \in \mathbb{R}^{p \times p}$ only appears on the GCN of subject k . Thus, it represents the individual connections for subject k not explained by moderator data. We illustrate the role of the additional fully individual structure in the GCNs across subjects in Figure 2.

To simplify using the notation introduced in (1), we write,

$$\mathbf{Y}_t^{(k)} = (\Phi^*(\mathbf{X}_k) + \Delta^{(Ik)*}) \mathbf{Y}_{t-1}^{(k)} + \boldsymbol{\varepsilon}_t^{(k)}, \quad \text{for all } 1 \leq k \leq n \text{ and } 2 \leq t \leq T_k. \quad (4)$$

We illustrate the model introduced in (4) in Figure 2.

Our goal is to estimate the common component $\Phi^{(0)*}$, GCN moderator effects $\{\Phi^{(M\ell)*}\}_{\ell=1}^p$, and fully individual GCN connections $\{\Delta^{(Ik)*}\}_{k=1}^n$.

3 Methodology

We introduce our method for simultaneously estimating the shared GCN connections across all subjects $\Phi^{(0)*}$, as well as the moderated effects $\{\Phi^{(M\ell)*}\}_{\ell=1}^p$ and fully individual connec-




Figure 2: Illustration of the MOD-VAR framework with $p = 2$ binary moderators and fully individual terms $\{\Delta^{(Ik)}\}_{k \in [n]}$. This models scenarios where fully individual connections may be present across GCNs, in addition to shared and moderator GCN connections.

tions $\{\Delta^{(Ik)*}\}_{k=1}^n$. We achieve this via an ℓ_1 -penalized time series estimation approach, which takes into account the estimation of all components. We first introduce the method when no individual remainders are present, *i.e.*, $\Delta^{(I1)*} = \Delta^{(I1)*} = \dots = \Delta^{(In)*} = \mathbf{0}$.

3.1 Mod-VAR With Only Moderator Components

Consider the case in which all inter-subject variability can be explained by our moderator variables, as described in (1). Then, our GCN follows the structure (4). Consider $\Phi = [\Phi^{(0)}, \Phi^{(M1)}, \Phi^{(M2)}, \dots, \Phi^{(Mp)}] \in \mathbb{R}^{d \times (d+dp)}$, and, for any vector of p moderators $\mathbf{x}$, we define $\Phi(\mathbf{x})$ as the $d \times d$ matrix given by

$$\Phi(\mathbf{x}) := \Phi^{(0)} + x_1 \Phi^{(M1)} + x_2 \Phi^{(M2)} + \dots + x_p \Phi^{(Mp)}. \quad (5)$$

Our goal is to estimate $\hat{\Phi}$. Since this requires the estimation of a total of $q = (p+1)d^2$ parameters, the estimation can be challenging, specially in cases where the number of study subjects n or number of time points per subject $\{T_k\}_{k=1}^n$ are limited. In such scenarios, one may estimate Φ^* under the assumption of sparsity, *i.e.*, only a small number of parameters in Φ^* are non-zero, with the vast majority of the entries in Φ^* being zero. As seen in the literature, the estimation of such sparse VAR structures can be achieved through the use of ℓ_1 -penalties [1; 6]. Here, we adapt such sparse-estimation techniques for detecting

moderator effects in the VAR models of interest.

To derive a sparse estimation of Φ^* , we introduce the MOD-VAR optimization problem

$$\hat{\Phi} = \underset{[\Phi^{(0)}, \Phi^{(M1)}, \dots, \Phi^{(Mp)}]}{\operatorname{argmin}} \frac{1}{2N} \sum_{k=1}^n \sum_{t=2}^{T_k} \|\mathbf{Y}_t^{(k)} - \Phi(\mathbf{X}_k) \mathbf{Y}_{t-1}^{(k)}\|_2^2 + \lambda_1 \|\Phi^{(0)}\|_1 + \lambda_2 \sum_{\ell=1}^p \|\Phi^{(M\ell)}\|_1. \quad (6)$$

where $N = T_1 + T_2 + \dots + T_n - n$, and $\lambda_1, \lambda_2 > 0$ are tuning parameters. Observe that the optimization (6) aims to minimize the mean square forecasting error (MSFE) across all subjects. Furthermore, it incorporates a non-uniform penalty, that induces different levels of sparsity for the shared component $\Phi^{(0)}$ and the moderator effects $\{\Phi^{(M\ell)}\}_k$.

We can simplify the expression of (6) to a reduced matrix form. For this, consider $\mathbf{Z}^{(k)} = [\mathbf{Y}_T^{(k)}, \mathbf{Y}_{T-1}^{(k)}, \dots, \mathbf{Y}_2^{(k)}]'$, $\mathbf{Y}^{(k)} = [\mathbf{Y}_{T-1}^{(k)}, \mathbf{Y}_{T-2}^{(k)}, \dots, \mathbf{Y}_1^{(k)}]'$, and $\mathbf{E}^{(k)} = [\boldsymbol{\varepsilon}_T^{(k)}, \boldsymbol{\varepsilon}_{T-1}^{(k)}, \dots, \boldsymbol{\varepsilon}_2^{(k)}]'$. Here the matrices $\mathbf{Z}^{(k)}$, $\mathbf{Y}^{(k)}$ and $\mathbf{E}^{(k)}$ aggregate the contemporaneous, lagged and error terms for subject k , respectively. Furthermore, we introduce

$$\mathbf{W}^{(k)} := [\mathbf{Y}^{(k)}, X_{k1} \mathbf{Y}^{(k)}, \dots, X_{kp} \mathbf{Y}^{(k)}] = \mathbf{X}_k' \otimes \mathbf{Y}^{(k)} \in \mathbb{R}^{(T_k-1) \times (pd+d)}.$$

One can show that $(\mathbf{Y}^{(k)})' \cdot \Phi^*(\mathbf{X}_k) = (\mathbf{W}^{(k)}) \cdot (\Phi^*)'$. We gather data from all subjects through the matrices $\mathbf{Z} := [\mathbf{Z}^{(1)'}, \mathbf{Z}^{(2)'}, \dots, \mathbf{Z}^{(n)'}]'$, $\mathbf{W} := [\mathbf{W}^{(1)'}, \mathbf{W}^{(2)'}, \dots, \mathbf{W}^{(n)'}]'$, and $\mathbf{E} := [\mathbf{E}^{(1)'}, \mathbf{E}^{(2)'}, \dots, \mathbf{E}^{(n)'}]'$. We can summarize all relationships in (1) with the simplified matrix notation $\mathbf{Z} = \mathbf{W} \cdot (\Phi^*)' + \mathbf{E}$. The optimization problem (6) is equivalent to solving

$$\hat{\Phi} := \underset{\Phi=[\Phi^{(0)}, \Phi^{(M1)}, \dots, \Phi^{(Mp)}]}{\operatorname{argmin}} \frac{1}{2N} \|\mathbf{Z} - \mathbf{W}\Phi'\|_F^2 + \lambda_1 \|\Phi^{(0)}\|_1 + \lambda_2 \|\Phi^{(M1)}, \dots, \Phi^{(Mp)}\|_1. \quad (7)$$

Thus, the estimation of $\hat{\Phi}$ can be reduced to a multiple regression optimization problem, with a non-uniform ℓ_1 penalty. We describe the optimization algorithm in Section 4.

3.2 Mod-VAR With Fully Idiographic Components

In many real data scenarios, moderator data may not be available, or the heterogeneity across subject may not be fully explained by the available moderator data. In such cases, we may aim to estimate additional fully individual connections in the GCN of each subject. We construct a set of artificial moderators to adapt optimization problem described in Section 3.1 for estimating common and fully individual components of the GCNs.

We consider the case in which we do not have moderator data available, and simply aim to estimate the shared and individual structure across multiple GCNs. This scenario is equivalent to the problem studied by the multi-VAR method [6]. For each subject k , we

define the artificial moderator vector $\mathbf{X}_k^{(I)} = (0, \dots, 0, 1, 0, \dots, 0)'$, a column vector of size $n \times 1$, with zero in all entries except the k -th entry. Since $\mathbf{X}_k^{(I)}$ has n artificial moderators, the associated parameters of MOD-VAR model can be expressed in the form of a $d \times (nd + d)$ matrix $\Phi = [\Phi^{(0)}, \Phi^{(M1)}, \Phi^{(M2)}, \dots, \Phi^{(Mn)}]$. Observe that, since $\mathbf{X}_k^{(I)}$ is indicator of the k -th entry and (5), $\Phi(\mathbf{X}_k^{(I)}) = \Phi^{(0)} + \Phi^{(Mk)}$, for all $k \in [n]$. Since the moderator effect $\Phi^{(Mk)}$ appears only in the coefficients of subject k , $\Phi^{(Mk)}$ corresponds to a fully individual connections in the GCN for subject k . From this, we change our notation for the estimating parameter to $\Phi = [\Phi^{(0)}, \Delta^{(I1)}, \Delta^{(I2)}, \dots, \Delta^{(In)}]$, where $\Delta^{(Ik)}$ represents the fully individual GCN connections to subject k . In summary, the estimation of $\Phi = [\Phi^{(0)}, \Delta^{(I1)}, \Delta^{(I2)}, \dots, \Delta^{(In)}]$ can be achieved by solving our proposed MOD-VAR optimization problem (7) with the artificial moderator design matrix $\mathbf{X}_{(I)} := [\mathbf{X}_1^{(I)}, \mathbf{X}_2^{(I)}, \dots, \mathbf{X}_n^{(I)}]' = I_n \in \mathbb{R}^{n \times n}$.

We may also be interested in scenarios where the GCNs incorporates shared, moderator-based and fully individual connections across subjects, as described in Section 2.3. For this, consider, for n subjects, the $n \times p$ moderator data matrix $\mathbf{X}_{(M)} = [\mathbf{X}_1, \mathbf{X}_2, \dots, \mathbf{X}_n]' \in \mathbb{R}^{n \times p}$, and the artificial moderator matrix $\mathbf{X}_{(I)} \in \mathbb{R}^{n \times n}$. We construct the expanded matrix $\mathbf{X}_{(M,I)} := [\mathbf{X}_{(M)}, \mathbf{X}_{(I)}]$. By incorporating both data-driven and artificial moderators to our matrix $\mathbf{X}_{(M,I)}$, we can estimate all the parameters $\Phi^{(0)}$, $\{\Phi^{(M\ell)}\}_{\ell=1}^p$ and $\{\Delta^{(Ik)}\}_{k=1}^n$ simultaneously by solving our proposed optimization problem (7).

4 Computational Algorithm and Estimation

We now discuss the procedure for optimizing (7), as well as for selecting the optimal tuning parameters λ_1 and λ_2 for fitting the MOD-VAR method.

4.1 Column-wise FISTA

The estimation of MOD-VAR requires solving the optimization problem (7). Since both terms $\|\mathbf{Z} - \mathbf{W}\Phi'\|_F^2/2N$ and $\mathcal{P}(\Phi) = \lambda_1 \|\Phi^{(0)}\|_1 + \lambda_2 \|\Phi^{(M1)}, \dots, \Phi^{(Mp)}\|_1$ are decomposable by the rows of the matrix Φ , we can solve for Φ by solving for each row separately. Thus, for each $\ell \in [d]$ we aim to solve the optimization problem

$$\hat{\phi}_\ell = \underset{\phi \in \mathbb{R}^{d+dp}}{\operatorname{argmin}} \left\{ \frac{1}{2N} \|\mathbf{Z}_\ell - \mathbf{W}\phi\|_2^2 + \lambda_1 \|\phi_{[1:d]}\|_1 + \lambda_2 \|\phi_{[(d+1):(d+dp)]}\|_1 \right\}. \quad (8)$$

Once we obtain $\hat{\phi}_1, \hat{\phi}_2, \dots, \hat{\phi}_d \in \mathbb{R}^d$, we construct $\hat{\Phi} = [\hat{\phi}_1, \hat{\phi}_2, \dots, \hat{\phi}_d]' \in \mathbb{R}^{d \times (d+dp)}$. Then, the estimated common portion of the GCN is defined as $\hat{\Phi}^{(0)} = \hat{\Phi}_{\cdot, [1:d]}$, and, for any $k \in [p]$ the estimated effect of the k -th moderator on the GCN defined as $\hat{\Phi}^{(Mk)} = \hat{\Phi}_{\cdot, [(d+d(k-1)):(d+dk)]}$. A summary of this procedure is provided in Algorithm 1. As discussed in Section 3.2, fully individual components can be incorporated into the model via artificial moderator variables. Therefore, the procedure described in Algorithm 1 can be applied for estimating shared, moderated and fully individual connections across GCNs, by simple modifications of the moderator design matrix $\mathbf{X}$.

Algorithm 1: MOD-VAR.

Input: Time series $\{\mathbf{Y}_t^{(k)} : k \in [n], t \in [T_k]\} \subset \mathbb{R}^d$, moderators $\mathbf{X} \in \mathbb{R}^{n \times p}$, tuning parameters $\lambda_1, \lambda_2 > 0$.

for $k \in [n]$ **do**

- Define $\mathbf{Z}^{(k)} = [\mathbf{Y}_2^{(k)}, \mathbf{Y}_3^{(k)}, \dots, \mathbf{Y}_{T_k}^{(k)}]' \in \mathbb{R}^{(T_k-1) \times d}$.
- Define $\mathbf{Y}^{(k)} = [\mathbf{Y}_1^{(k)}, \mathbf{Y}_2^{(k)}, \dots, \mathbf{Y}_{T_k-1}^{(k)}]' \in \mathbb{R}^{(T_k-1) \times d}$.
- Set $\mathbf{X}_k^* := (1, \mathbf{X}_{k\cdot}')' \in \mathbb{R}^{p+1}$.
- Define $\mathbf{W}^{(k)} = \mathbf{X}_k^* \otimes \mathbf{Y}^{(k)} = [\mathbf{Y}^{(k)}, X_{k1}\mathbf{Y}^{(k)}, \dots, X_{kp}\mathbf{Y}^{(k)}] \in \mathbb{R}^{(T_k-1) \times (d+dp)}$.

end

Define $\mathbf{Z} = [\mathbf{Z}^{(1)'}, \mathbf{Z}^{(2)'}, \dots, \mathbf{Z}^{(n)'}]' \in \mathbb{R}^{N \times d}$.

Define $\mathbf{W} = [\mathbf{W}^{(1)'}, \mathbf{W}^{(2)'}, \dots, \mathbf{W}^{(n)'}]' \in \mathbb{R}^{N \times (d+dp)}$.

for $\ell \in [d]$ **do**

- Solve $\hat{\phi}_\ell = \operatorname{argmin}_{\phi \in \mathbb{R}^{d(p+1)}} \|\mathbf{Z}_\ell - \mathbf{W}\phi\|_2^2/2N + \lambda_1 \|\phi_{1:d}\|_1 + \lambda_2 \|\phi_{(d+1):(d+pd)}\|_1$.

end

Define $\hat{\Phi} := \hat{\Phi}(\lambda_1, \lambda_2) = [\hat{\phi}_1, \hat{\phi}_2, \dots, \hat{\phi}_d]' \in \mathbb{R}^{d \times (d+pd)}$.

Set $\hat{\Phi}^{(0)} := \hat{\Phi}^{(0)}(\lambda_1, \lambda_2) = \hat{\Phi}_{\cdot, [1:d]}$.

For $s \in [p]$, define $\hat{\Phi}^{(Ms)} := \hat{\Phi}^{(Ms)}(\lambda_1, \lambda_2) = \hat{\Phi}_{\cdot, [(d+d(s-1)):(d+ds)]}$.

Output: Common effects $\hat{\Phi}^{(0)}$, and moderated effects $\{\hat{\Phi}^{(Ms)}\}_{s=1}^p$

In order to solve the optimization problem (8) for each $\ell \in [d]$, we apply the fast iterative shrinkage-threshold algorithm (FISTA) algorithm [3]. For each $\ell \in [d]$, we define $\mathcal{L}^{(\ell)}(\phi) = -\phi' \mathbf{W}' \mathbf{W} \phi / 2N - \mathbf{Z}_\ell' \mathbf{W} \phi / N$, and the penalty term $\mathcal{P}(\phi) = \lambda_1 \|\phi_{1:d}\|_1 + \lambda_2 \|\phi_{(d+1):(d+pd)}\|_1$. It can be shown that optimizing (8) is equivalent to solving the optimization problem $\operatorname{argmin}_\phi \mathcal{L}^{(\ell)}(\phi) + \mathcal{P}(\phi)$. Furthermore, FISTA requires calculating a Lipschitz constant $L > 0$ for the gradient $\nabla \mathcal{L}^{(\ell)}(\phi) = (\mathbf{W}' \mathbf{W} / N) \phi - \mathbf{W}' \mathbf{Z}_\ell / N$. The inputs of our implementation of the FISTA algorithm applied to each $\ell \in [d]$ are derived from the matrices $\mathbf{M} := \mathbf{W}' \mathbf{W} / N$, $\mathbf{b}_\ell := \mathbf{W}' \mathbf{Z}_\ell$, the tuning parameters $\lambda_1, \lambda_2 > 0$, the Lipschitz constant $L := \|\mathbf{W}' \mathbf{W}\|_2 / 2N$,

and a maximum iteration number I , and a tolerance level $\varepsilon > 0$. In addition, we may improve efficiency of our estimation by utilizing a warm start $\mathbf{x}_0 \in \mathbb{R}^{d+dp}$.

Our FISTA implementation consists of the following steps. We start by setting $s_1 = 1$, and initializing $\mathbf{x}_0$ if a warm start is not provided. Then, we set $\mathbf{y}_1 = \mathbf{x}_0$, and define the vector of penalization weights as $\mathbf{w} = (\lambda_1 \mathbf{1}'_d, \lambda_2 \mathbf{1}'_{dp})' \in \mathbb{R}^{d+dp}$. Then, for any $i = 1, 2, \dots, I$, the updates for $\mathbf{x}_i$, s_i and $\mathbf{y}_{i+1}$ are given by the equations

$$\begin{aligned}\mathbf{x}_i &:= ST_{\mathbf{w}/L}(\mathbf{y}_i - \nabla \mathcal{L}^{(\ell)}(\mathbf{y}_i)/L), \\ s_{i+1} &:= (1 + \sqrt{1 + 4s_i^2})/2, \\ \mathbf{y}_{i+1} &:= \mathbf{x}_i + ((s_i + 1)/s_{i+1})(\mathbf{x}_i - \mathbf{x}_{i-1}),\end{aligned}$$

where, for any $\mathbf{w}, \mathbf{a} \in \mathbb{R}^m$, we define $ST_{\mathbf{w}}(\mathbf{a}) \in \mathbb{R}^m$ as the weighted entrywise soft thresholding function given by the entrywise value $[ST_{\mathbf{w}}(\mathbf{a})]_i = \text{sign}(a_i) \cdot \max\{0, |a_i| - |w_i|\}$. The full procedure is summarized in Algorithm 2.

Algorithm 2: FISTA for ℓ -step of MOD-VAR estimation.

Input: $\mathbf{M} := \mathbf{W}'\mathbf{W}/N \in \mathbb{R}^{(d+dp) \times (d+dp)}$, $\mathbf{b}_\ell := \mathbf{W}'\mathbf{Z}_\ell/N \in \mathbb{R}^{d+dp}$, hyperparameters $\lambda_1, \lambda_2 > 0$, Lipschitz constant $L := \|\mathbf{M}\|_2$, maximum iterations $I > 0$, tolerance $\varepsilon > 0$, warm start $\mathbf{x}_0 \in \mathbb{R}^{d+dp}$.

Initialize $s_1 = 1$, $\mathbf{y}_1 = \mathbf{x}_0$, and $\mathbf{x}_1 = ST_{\mathbf{w}/L}(\mathbf{y}_1 - \nabla \mathcal{L}^{(\ell)}(\mathbf{y}_1)/L)$.

Set $\mathbf{w} = (\lambda_1 \mathbf{1}'_d, \lambda_2 \mathbf{1}'_{dp})' \in \mathbb{R}^{d+dp}$ and $i = 1$.

while $i \leq I$ and $\|\mathbf{x}_i - \mathbf{x}_{i-1}\|_2 > \varepsilon$ **do**

 Calculate $\nabla \mathcal{L}^{(\ell)}(\mathbf{y}_i) = \mathbf{M}\mathbf{y}_i - \mathbf{b}_\ell$.

 Set $\mathbf{x}_i := ST_{\mathbf{w}/L}(\mathbf{y}_i - \nabla \mathcal{L}^{(\ell)}(\mathbf{y}_i)/L)$.

 Set $s_{i+1} = (1 + \sqrt{1 + 4s_i^2})/2$.

 Set $\mathbf{y}_{i+1} = \mathbf{x}_i + ((s_i + 1)/s_{i+1})(\mathbf{x}_i - \mathbf{x}_{i-1})$.

 Set $i \leftarrow i + 1$.

end

Set $\hat{\phi}_\ell := \hat{\phi}_\ell(\lambda_1, \lambda_2) = \mathbf{x}_{i-1}$.

4.2 Tuning Parameter Selection

To select the optimal MOD-VAR fit, we must select the optimal $\lambda_1, \lambda_2 > 0$ over a range of potential values. We describe three strategies for the selection of $\lambda_1, \lambda_2 > 0$: by-subject cross validation, rolling window and cross validation, and an adaptive version of our proposed moderator vector auto-regressive model (adaMOD-VAR).

4.2.1 By Subject Cross-Validation

We introduce the by-subject cross-validation (BSCV) procedure for fitting our proposed MOD-VAR method. Our J -fold BSCV procedure is as follows. First, we divide our set of subjects $[n]$ into J approximately equally sized partitions $\mathcal{C}_1, \mathcal{C}_2, \dots, \mathcal{C}_J$. Then, for each $j \in [J]$, and choice of tuning parameters $\lambda_1 > 0$, $\lambda_2 > 0$, we fit $\hat{\Phi}_{(BSCV,j)}(\lambda_1, \lambda_2)$ as the solution to (6) based exclusively on the data of the subjects $[n] \setminus \mathcal{C}_j$. This provides us with the common GCN structure $\hat{\Phi}_{(BSCV,j)}^{(0)}(\lambda_1, \lambda_2) \in \mathbb{R}^{d \times d}$ and for each $s \in [p]$, the effect of the s -moderator onto the GCN $\{\hat{\Phi}_{(BSCV,j)}^{(M\ell)}(\lambda_1, \lambda_2)\}_{\ell=1}^p \subset \mathbb{R}^{d \times d}$. Then, for each subject $k \in \mathcal{C}_j$, we define the cross-validation predicted GCN based on the moderator vector $\mathbf{X}_k$,

$$\hat{\Phi}_{(BSCV,j)}^{(\lambda_1, \lambda_2)}(\mathbf{X}_k) = \hat{\Phi}_{(BSCV,j)}^{(0)}(\lambda_1, \lambda_2) + \sum_{s=1}^p X_{ks} \hat{\Phi}_{(BSCV,j)}^{(Ms)}(\lambda_1, \lambda_2) \in \mathbb{R}^{d \times d}.$$

Here, $\hat{\Phi}_{(BSCV,j)}^{(\lambda_1, \lambda_2)}(\mathbf{X}_k)$ provides an estimation of the GCN corresponding to subject $k \in \mathcal{C}_j$, with no use of the data of subject k , instead relying only on the for the subject set $[n] \setminus \mathcal{C}_j$.

Finally, we define the by-subject cross-validation mean-squared forecasting error as

$$\begin{aligned} BSCV.MSFE(\lambda_1, \lambda_2) &= \frac{1}{J} \sum_{j \in [J]} \sum_{k \in \mathcal{C}_j} \frac{1}{T_k - 1} \left\| \mathbf{Z}^{(k)} - \mathbf{Y}^{(k)} \hat{\Phi}_{(BSCV,j)}^{(\lambda_1, \lambda_2)}(\mathbf{X}_k) \right\|_F^2 \\ &= \frac{1}{J} \sum_{j \in [J]} \sum_{k \in \mathcal{C}_j} \frac{1}{T_k - 1} \sum_{t=2}^{T_k} \left\| \mathbf{Y}_t^{(k)} - \hat{\Phi}_{(BSCV,j)}^{(\lambda_1, \lambda_2)}(\mathbf{X}_k) \mathbf{Y}_{t-1}^{(k)} \right\|_2^2 \end{aligned}$$

We select the pair $\lambda_1, \lambda_2 > 0$ that minimize the $BSCV.MSFE(\lambda_1, \lambda_2)$.

This approach enjoys major computational advantages compared to other procedures like rolling-window cross-validation, since it only requires the fit of J runs of our MOD-VAR fitting procedure. We emphasize that by-subject cross validation can be used exclusively when no fully individual connections are expected across models.

4.2.2 Rolling Window Cross-Validation

Rolling window cross-validation is an alternative procedure for the selection of an optimal MOD-VAR model, based on the method proposed in [19]. First, we select $T = \min_{k \in [n]} \{T_k - 1\}$, and we set $J = \lfloor T/3 \rfloor$. Then, for any $j \in [J]$, we select the observations $\cup_{k \in [n]} \{Y_t^{(k)} : t \in [T_k - j - 2]\}$ as the training set. With this, we obtain the by-subject GCN estimators $\hat{\Phi}_{(RWCV,j)}^{(k; \lambda_1, \lambda_2)} \in \mathbb{R}^{d \times d}$ for each potential choice of $\lambda_1, \lambda_2 > 0$. Finally, we measure the rolling

window cross validation mean squared forecasting error, given by,

$$RWCV.MSFE(\lambda_1, \lambda_2) = \frac{1}{J} \sum_{j \in [J]} \frac{1}{n} \sum_{k \in [n]} \|\mathbf{Y}_{T_k-j-1}^{(k)} - \hat{\Phi}_{(RWCV,j)}^{(k;\lambda_1,\lambda_2)} \cdot \mathbf{Y}_{T_k-j}^{(k)}\|_2^2.$$

We select the values $\lambda_1^*, \lambda_2^* > 0$ with the lowest RWCV.MSFE value.

4.2.3 Adaptative MOD-VAR

To improve upon the performance of our original estimator, we propose a refinement of our MOD-VAR estimator through an adaptive approach [30]. For this, we consider $\hat{\Phi} = [\hat{\Phi}^{(0)}, \hat{\Phi}^{(M1)}, \dots, \hat{\Phi}^{(Mp)}, \hat{\Delta}^{(I1)}, \dots, \hat{\Delta}^{(In)}] \in \mathbb{R}^{d \times (d+dp+dn)}$ be an initial estimator of our proposed MOD-VAR. Then, given a constant $\gamma > 0$ we define the adaptive weight for the shared and moderated coefficients $[\hat{\Phi}^{(0)}, \hat{\Phi}^{(M1)}, \dots, \hat{\Phi}^{(Mp)}]$ as,

$$\omega_{ada}(\hat{\Phi})_{ij} = |\hat{\Phi}_{ij}|^{-\gamma}, \quad \forall 1 \leq i \leq d, 1 \leq j \leq p+1.$$

For the fully individual connectivity portion $[\hat{\Delta}^{(I1)}, \dots, \hat{\Delta}^{(In)}]$, we select our adaptive penalty similar to the proposal of Fisher et al. [7]. First, based on $\hat{\Phi}$, we calculate the by-subject GCNs, $\hat{\Phi}^{(1)}, \hat{\Phi}^{(2)}, \dots, \hat{\Phi}^{(n)} \in \mathbb{R}^{d \times d}$. Then, we calculate the entrywise median GCN $\tilde{\Phi} \in \mathbb{R}^{d \times d}$, where $\tilde{\Phi}_{ij}$ is the median of $\hat{\Phi}_{ij}^{(1)}, \hat{\Phi}_{ij}^{(2)}, \dots, \hat{\Phi}_{ij}^{(n)}$. Finally, we define the penalty adaptive weight for $\Delta_{ij}^{(Ik)}$ as,

$$\omega_{ada}(\hat{\Phi})_{i,d(p+k)+j} = |\Delta_{ij}^{(Ik)} - \tilde{\Phi}_{ij}|^{-\gamma}, \quad \forall 1 \leq i, j \leq d, 1 \leq k \leq n.$$

From this, we define the adaptively-weighted ℓ_1 penalty for Φ as,

$$P_{ada}(\Phi) = \sum_{\ell=1}^p \sum_{i,j=1}^d \omega_{ada}(\hat{\Phi})_{ij} \cdot |\Phi_{ij}|. \quad (9)$$

Then, our adaptive MOD-VAR estimator is given as a solution to the optimization problem

$$\hat{\Phi}_{ada}(\lambda) = \underset{\Phi \in \mathbb{R}^{d \times (d+dp)}}{\operatorname{argmin}} \frac{1}{2N} \|\mathbf{Z} - \mathbf{W}\Phi\|_F^2 + \lambda P_{ada}(\Phi), \quad \text{for } \lambda > 0.$$

As observed in our numerical simulations, the adaptive procedure has significant improvements in terms of both forecasting, consistency, and edge-recovery performance.

5 Theoretical Properties

We explore the theoretical properties of our proposed MOD-VAR, and provide high-dimensional consistency guarantees of estimation. In particular, we are interested in the estimation error $\|\hat{\Phi} - \Phi^*\|_F$ and ℓ_1 -error $\|\hat{\Phi} - \Phi^*\|_1$. To explore this, consider the MOD-VAR matrix

of true parameters Φ^* is s -sparse: $\|\Phi^*\|_0 \leq s$. To simplify our exposition, we assume that all time series have the same number of time observations, *i.e.*, $T := T_1 = T_2 = \dots = T_n$. Then, we denote N the total number of observations available as $N = n(T - 1)$.

To establish our guarantees, consider the following set of assumptions.

Assumption 1 (Bounded moderators). *Assume the moderator vectors $\mathbf{X}, \mathbf{X}_1, \mathbf{X}_2, \dots, \mathbf{X}_n$ are independent, identically distributed and entrywise bounded random vectors *i. e.*, $|X_i| \leq 1$, with mean zero, and covariance matrix $\Sigma_X = \mathbb{E}[\mathbf{X}\mathbf{X}']$. Furthermore, assume that for some constant $0 < \eta < 1$, we have that $\sup_{\|\mathbf{x}\|_\infty \leq 1} \|\Phi^*(\mathbf{x})\|_2 < \eta < 1$*

Assumption 2 (Curvature condition). *Let $\Lambda_Y := \inf_{\|\mathbf{x}\|_\infty \leq 1} \Lambda_{\min}(\text{Cov}(Y_t^k | X_k))$, and $\alpha = \Lambda_Y \cdot \Lambda_{\min}(\Sigma_X)/2$. There exists a universal constant $C > 0$ such that, for some $t > 0$,*

$$\alpha > Cs^2 \left(\sqrt{\frac{\log(d^2p) + t}{n}} + \frac{\log(2n)(\log(pd^2) + t)}{n} \right).$$

Our Assumption 1 requires that the moderator variables satisfy the entrywise bound $|X_i| \leq 1$. This is not restrictive, since in applications, covariates often correspond to categorical or bounded score variables. For example, in the analysis of emotional/behavioral data from intensive longitudinal studies, moderators of interest are often gender, socio-economic status, or ethnicity age, which can be coded and normalized as discrete, centered and bounded variables. The condition $\sup_{\|\mathbf{x}\|_\infty \leq 1} \|\Phi^*(\mathbf{x})\|_2 < \eta < 1$ serves to ensure that all the sampled time series are stable [1].

The convergence of our MOD-VAR estimator relies on ensuring sufficient curvature of the matrix $\mathbf{W}'\mathbf{W}/N \in \mathbb{R}^{(d+dp) \times (d+dp)}$. Since the design matrix $\mathbf{W}$ is composed of the blocks $\mathbf{X}_k^* \otimes \mathbf{Y}^{(k)}$, where $\mathbf{X}_k^* = (1, \mathbf{X}_k')' \in \mathbb{R}^{p+1}$, the minimum eigenvalue of $\mathbb{E}[\mathbf{W}'\mathbf{W}/N]$ can be controlled by imposing minimum eigenvalue conditions of Σ_X and the conditional covariances of $\text{Cov}(\mathbf{Y}_t^{(k)} | \mathbf{X}_k)$. Similar minimum eigenvalue controls are commonly used in the literature of high-dimensional regression literature [16; 1; 25].

Under Assumptions 1 and 2, we derive high-dimensional guarantees of convergence.

Theorem 1. *Consider multiple subject data $\{(\mathbf{Y}_k, \mathbf{X}_k)\}_{k=1}^n$ that follows the MOD-VAR model (4). Furthermore, suppose Assumption 1 holds, Assumption 2 is satisfied for a given $t > 0$, and $n > 16 \log(2n) \log(pd^2)$. Let $C_M := \Lambda_{\max}(\Sigma_\varepsilon)(2 + 4\eta/(1 - \eta)^2)$. Then, assuming that $\lambda := \lambda_1 = \lambda_2$ satisfies the inequality $\lambda > 16\pi C_M \sqrt{\log(pd^2)/N}$, we have that, with*

probability at least $1 - 3e^{-t} - (6/(pd^2))^2$,

$$\|\hat{\Phi} - \Phi^*\|_2 \leq \frac{12\sqrt{s}\pi C_M}{\alpha} \sqrt{\frac{\log(pd^2)}{N}}; \quad (10)$$

$$\|\hat{\Phi} - \Phi^*\|_1 \leq \frac{72s\pi C_M}{\alpha} \sqrt{\frac{\log(pd^2)}{N}}. \quad (11)$$

Our Theorem 1 provides guarantees on the convergence of our proposed MOD-VAR, even in high dimensional scenarios. First, notice that our bound on $\|\hat{\Phi} - \Phi^*\|_2$ can be translated into a bound on the individual GCN estimation,

$$\begin{aligned} \max_{k=1,\dots,n} \|\hat{\Phi}(\mathbf{X}_k) - \Phi^*(\mathbf{X}_k)\|_1 &\leq \|\hat{\Phi}^{(0)} - \hat{\Phi}^{(0)*}\|_1 + \max_{k=1,\dots,n} \sum_{i=1}^p |X_{ki}| \|\hat{\Phi}^{(Mi)} - \hat{\Phi}^{(Mi)*}\|_1 \\ &\leq \|\hat{\Phi} - \Phi^*\|_1 \end{aligned}$$

As a consequence of this, as long as $\log(pd^2) = O(N)$, we may ensure the ℓ_1 -norm consistency of our proposed MOD-VAR estimator. The rate $O(\sqrt{(\log(p) + 2\log(d))/N})$ described in (11) depends on the total sample size $N = \sum_{i=1}^n (T_i - 1)$. As long as $\log(p)$ is moderate (for example, if $p = O(d)$), the rate obtained in Theorem 1 represents a significant improvement to the rate of $O(\sqrt{2\log(d)/T})$ attained by the individual VAR estimation without moderators. This demonstrates the significant improvement of performing joint estimation when moderator effects are present in our time series model. Furthermore, notice that since the convergence depends on N and not on T , consistent estimation is possible, even in cases where the number of available observations per subject T is limited, as long as n the number of subjects n is large. This is possible thanks to the knowledge transfer across GCN's facilitated by the moderator data. The main requirement on the number of subjects n is $\log(pd^2) = o(n/\log(n))$, found in Assumption 2. This ensures that there is sufficient heterogeneity in the moderator data, and consistent estimation of Σ_X , which in turn facilitates the estimation of the GCN moderator effects.

6 Numerical Simulations

To assess the performance of our proposed MOD-VAR, and compare it to methods in the literature, we perform a set of numerical simulation experiments.

6.1 Data Generating Method

To generate our data, we select a dimension $d \in \{10, 20\}$, number of available moderator variables $p \in \{2, 5\}$, number of available time series $n \in \{15, 20, 25\}$, and number of available observations per time series $T \in \{30, 60, 90\}$.

To generate our data, we generate a matrix of moderator variables $\mathbf{X} = [\mathbf{X}_1, \mathbf{X}_2, \dots, \mathbf{X}_n] \in \mathbb{R}^{n \times p}$, as well as matrices of shared GCN connections $\Phi^{(0)*} \in \mathbb{R}^{d \times d}$, the GCN moderator effects $\{\Phi^{(Ms)*}\}_{s=1}^p \subset \mathbb{R}^{d \times d}$, and the matrices of fully individual connections $\{\Delta^{(Ik)*}\}_{k=1}^n \subset \mathbb{R}^{d \times d}$. Once these matrices are generated, for each $k \in [n]$, we define the GCN for the k -th time series based on the moderators $\mathbf{X}_k = (X_{k1}, \dots, X_{kp})'$ as,

$$\Phi^{(k)*} = \Phi^{(0)*} + \sum_{s=1}^p X_{ks} \Phi^{(Ms)*} + \Delta^{(Ik)*} \in \mathbb{R}^{d \times d}.$$

Then, for each $k \in [n]$, we generate the time series data $\{\mathbf{Y}_t^{(k)}\}_{t=1}^{T+400}$, by first setting $\mathbf{Y}_0^{(k)} = \mathbf{0}$, and $\mathbf{Y}_t^{(k)} = \Phi^{(k)*} \mathbf{Y}_{t-1}^{(k)} + \boldsymbol{\varepsilon}_t$, where $\{\boldsymbol{\varepsilon}_t\}_{t=1}^{T+403} \sim N_d(\mathbf{0}, (0.1) \cdot \mathbf{I}_d)$. Then, we discard the first 400 time points of each time series. The additional 3 time points will be used to evaluate out-of-sample forecasting performance.

The simulation of our data requires a procedure for generating the matrix of moderator variables $\mathbf{X} = [\mathbf{X}_1, \mathbf{X}_2, \dots, \mathbf{X}_n] \in \mathbb{R}^{n \times p}$, and the GNC matrices $\Phi^{(0)*}$, $\{\Phi^{(Ms)*}\}_{s=1}^p$ and $\{\Delta^{(Ik)*}\}_{k=1}^n$. To generate $\mathbf{X} = [\mathbf{X}_1, \mathbf{X}_2, \dots, \mathbf{X}_n]' \in \mathbb{R}^{n \times p}$, we set $\mathbf{X}_k = (X_{k1}, \dots, X_{kp}) \stackrel{i.i.d.}{\sim} (0.25) \cdot \delta_0 + (0.75) \cdot \text{Unif}([-1, -0.5] \cup [0.5, 1])$. To generate the GNC matrices, we introduce $q_0, q_m, q_\Delta \in [0, 1]$ the probabilities of shared, moderated and fully individual connection probabilities. With these parameters in place, we generate our data as follows. First, we generate the shared network $\Phi^{(0)*} \in \mathbb{R}^{d \times d}$, moderated networks $\{\Phi^{(Ms)*}\}_{s=1}^p \subset \mathbb{R}^{d \times d}$, and fully individual network connections $\{\Delta^{(Ik)*}\}_{k=1}^n \subset \mathbb{R}^{d \times d}$ as follows:

$$\begin{aligned} \{\Phi_{ij}^{(0)*}\}_{i,j=1}^d &\stackrel{i.i.d.}{\sim} (1 - q_0)\delta_0 + q_0 \text{Unif}([-0.9, -0.5] \cup [0.5, 0.9]); \\ \{\Phi_{ij}^{(Ms)*}\}_{i,j=1}^d &\stackrel{i.i.d.}{\sim} (1 - q_m)\delta_0 + q_m \text{Unif}([-0.9, -0.5] \cup [0.5, 0.9]); \\ \{\Delta_{ij}^{(Ik)*}\}_{i,j=1}^d &\stackrel{i.i.d.}{\sim} (1 - q_\Delta)\delta_0 + q_\Delta \text{Unif}([-0.9, -0.5] \cup [0.5, 0.9]). \end{aligned}$$

Thus, the parameters q_0 , q_m and q_Δ control the level of sparsity of the matrices $\Phi^{(0)*}$, $\{\Phi^{(Ms)*}\}_{s=1}^p$ and $\{\Delta^{(Ik)*}\}_{k=1}^n$. In particular, notice that if $q_m = 0$, it follows that $\Phi^{(Ms)*} = \mathbf{0}$ for all $s \in [p]$. Thus, $q_m = 0$ corresponds to a data generating scenario in which the moderator variables do not have any effect on the connectivity pattern across the GCNs.

Furthermore, $q_\Delta = 0$ implies $\Delta^{(I^k)} = \mathbf{0}$ for all $k \in [n]$, which corresponds to an absence of fully individual structure in our time series. In our simulations, we explore different choices of q_0 , q_m and q_Δ to explore the impact that moderator effects, fully individual effects and the absence of each have on the GNC estimation accuracy.

6.2 Methods and Performance Metrics

We compare our proposed MOD-VAR to methods in the literature of multiple time series modeling. First, we consider the vector auto-regressive (VAR) model [1] applied separately to each time series. This method is implemented via the `bigVAR` R package [18]. We also compare to the original and adaptive multi-VAR methods [6], implemented via the `multivar` package. Finally, we compare our MOD-VAR method, fitted with rolling-window cross-validation, BIC and adaptive procedures.

To evaluate the performance, we consider the mean ℓ_1 and Frobenius GNC estimation errors, and root mean squared forecasting errors, given by

$$\begin{aligned} E_{\ell_1}(\hat{\Phi}, \Phi^*) &= \frac{1}{n} \sum_{k=1}^n \|\hat{\Phi}^{(k)} - \Phi^{(k)*}\|_1, \\ E_F(\hat{\Phi}, \Phi^*) &= \frac{1}{n} \sum_{k=1}^n \left\| \hat{\Phi}^{(k)} - \Phi^{(k)*} \right\|_F, \\ MSFE(\hat{\Phi}, \Phi^*) &= \frac{1}{n} \left(\sum_{k=1}^n \frac{1}{3} \sum_{h=1}^3 \|\mathbf{Y}_{T+h-1}^{(k)} - \hat{\Phi}^{(k)} \mathbf{Y}_{T+h}^{(k)}\|_2^2 \right)^{1/2}, \end{aligned}$$

respectively. In our supplementary materials, we incorporate results on mean true positive rate of connection, mean false positive rate, and computational time.

6.3 Results: Only Moderator-Driven Heterogeneity

First, we explore the scenario in which all the heterogeneity across the time series models can be accounted by the p moderator variables. This corresponds to the case $q_\Delta = 0$, *i.e.*, $\Delta^{(I^k)} = 0$ for all $k \in [n]$. For this case, we explore the cases $(q_0, q_m) \in \{(0.075, 0.025), (0.025, 0.075)\}$. We provide visualizations for the E_{ℓ_1} , E_F and $MSFE$ for $d = 10$ in this section. We provide visualizations for the TPR, FPR in our Supplementary Materials.

As we observe in Figure 3, both the MOD-VAR and adaptive MOD-VAR attain the best performance in terms of ℓ_1 -norm error. Notice that, depending on the degree of com-




Figure 3: Visualization of the mean ℓ_1 GNC estimation error for our proposed MOD-VAR, as well as methods in the literature when the only source of GCN heterogeneity are explained by moderator variables. Our proposed MOD-VAR and adaptive MOD-VAR attain the best performance, for all choices of number of subjects n , moderators p and percentage of shared structure in the GCN.

monality among the GCNs, the second best estimation may be attained by the individual VAR estimation or the Multi-VAR method. In the scenario with the least cross-subject heterogeneity in which 75% of the GCN is shared, and $p = 2$, we have that the Multi-VAR method is the next best performing method. On the other hand, in scenarios with high heterogeneity where only 25% of entries are shared and $p = 5$, we have that the VAR method is the second best performing method. Our proposed MOD-VAR attains a significantly better performance, since it can model joint structure, and uses moderator data to drastically reduce the parameter complexity of estimating the heterogeneous structure.

6.4 Results: Moderator and Fully Individual Heterogeneity

To further explore the performance of our estimation procedure, accounting for the potential presence of both moderator effects, as well as fully individual connections in the GNC of some individuals, we perform an additional simulation scenario. For this, we set $(q_0, q_m) \in \{(0.075, 0.025), (0.025, 0.075)\}$ and $q_\Delta = 0.03$. We provide visualizations for the E_{ℓ_1} , E_F and $MSFE$ for $d = 10$ in this section. We provide visualizations for the TPR, FPR in our Supplementary Materials.




Figure 4: Visualization of the mean ℓ_1 GNC estimation error for our proposed MOD-VAR, as well as methods in the literature when heterogeneity in the GCNs is a combination of moderator effects and fully individual connections. Our proposed adaptive MOD-VAR attain the best performance, for all choices of number of subjects n , moderators p and percentage of shared structure in the GCN.

7 Analysis of Mood and Behavior Data

8 Discussion

In this work, we introduce the moderator vector auto-regressive model (MOD-VAR), for estimating the GCN associated to a set of heterogeneous multivariate time series data. Our proposed MOD-VAR can identify common structure across heterogeneous GCNs, and identify whether individual structure present in the model may be explained by external moderator variables. Our method outperforms methods in the literature in terms of both GNC recovery and forecasting error, as showcased by our extensive numerical experiments. Furthermore, our theoretical developments show that MOD-VAR attains fast rates of convergence, even in high-dimensional scenarios, as long as $n = \Omega(\log((n + p + 1)d^2))$.

This work opens a variety of new avenues of research. In particular, we consider important to further explore applications of our proposed MOD-VAR to other multi subject studies in psychology and neuroscience, similar to the works of [11; 21; 4]. Furthermore, our proposed method opens a variety of new avenues of methodological research. For example, it would be of interest to test the statistical significance of the effect of moderators on GCN

networks. Further extensions beyond the sparsity assumption are also of interest.

Acknowledgments

This work was supported by the UCR Latino and Latin-American Research Center writing group. Holly O’Rourke and Gene Brewer were supported by the Army Research Institute for Behavioral and Social Sciences grant W911NF-18-S-0005. José Á. Sánchez Gómez and Holly O’Rourke were supported by the UCR Research Seed Funding and Grant Program.

Supplementary Materials

We provide proofs, additional discussions and simulations in the Supplementary Materials.

References

- [1] Basu, S. and G. Michailidis (2015). Regularized estimation in sparse high-dimensional time series models. *The Annals of Statistics* 43(4), 1535–1567.
- [2] Basu, S., A. Shojaie, and G. Michailidis (2015). Network granger causality with inherent grouping structure. *The Journal of Machine Learning Research* 16(1), 417–453.
- [3] Beck, A. and M. Teboulle (2009). A fast iterative shrinkage-thresholding algorithm for linear inverse problems. *SIAM journal on imaging sciences* 2(1), 183–202.
- [4] Deshpande, G., S. LaConte, G. A. James, S. Peltier, and X. Hu (2009). Multivariate granger causality analysis of fmri data. *Human brain mapping* 30(4), 1361–1373.
- [5] Fisher, A. J., J. D. Medaglia, and B. F. Jeronimus (2018). Lack of group-to-individual generalizability is a threat to human subjects research. *Proceedings of the National Academy of Sciences* 115(27), E6106–E6115.
- [6] Fisher, Z. F., Y. Kim, B. L. Fredrickson, and V. Pipiras (2022). Penalized estimation and forecasting of multiple subject intensive longitudinal data. *psychometrika* 87(2), 403–431.
- [7] Fisher, Z. F., Y. Kim, V. Pipiras, C. Crawford, D. J. Petrie, M. D. Hunter, and C. F. Geier (2024). Structured estimation of heterogeneous time series. *Multivariate Behavioral Research* 59(6), 1270–1289.

- [8] Gates, K. M. and P. C. Molenaar (2012). Group search algorithm recovers effective connectivity maps for individuals in homogeneous and heterogeneous samples. *NeuroImage* 63(1), 310–319.
- [9] Hamilton, J. D. (2020). *Time series analysis*. Princeton university press.
- [10] Javanmard, A. and A. Montanari (2014). Confidence intervals and hypothesis testing for high-dimensional regression. *The Journal of Machine Learning Research* 15(1), 2869–2909.
- [11] Kim, J., W. Zhu, L. Chang, P. M. Bentler, and T. Ernst (2007). Unified structural equation modeling approach for the analysis of multisubject, multivariate functional mri data. *Human brain mapping* 28(2), 85–93.
- [12] Lane, S. T., K. M. Gates, H. K. Pike, A. M. Beltz, and A. G. Wright (2019). Uncovering general, shared, and unique temporal patterns in ambulatory assessment data. *Psychological methods* 24(1), 54.
- [13] Lee, C. and K. M. Gates (2026). Group iterative multiple model estimation approaches in clinical science. *Annual Review of Clinical Psychology* 22, 23–48.
- [14] Li, Y., J. Wood, L. Ji, S.-M. Chow, and Z. Oravecz (2022). Fitting multilevel vector autoregressive models in stan, jags, and mplus. *Structural equation modeling: a multidisciplinary journal* 29(3), 452–475.
- [15] Liu, R. and G. Yu (2024). A convex formulation of covariate-adjusted gaussian graphical models via natural parametrization. *arXiv e-prints*, arXiv–2410.
- [16] Loh, P.-L. and M. J. Wainwright (2012). High-dimensional regression with noisy and missing data: Provable guarantees with nonconvexity. *The Annals of Statistics* 40(3), 1637 – 1664.
- [17] Ni, Y., F. C. Stingo, and V. Baladandayuthapani (2022). Bayesian covariate-dependent gaussian graphical models with varying structure. *Journal of Machine Learning Research* 23(242), 1–29.

- [18] Nicholson, W., D. Matteson, J. Bien, and I. Wilms (2023). *BigVAR: Dimension Reduction Methods for Multivariate Time Series*. R package version 1.1.2.
- [19] Nicholson, W. B., D. S. Matteson, and J. Bien (2017). Varx-l: Structured regularization for large vector autoregressions with exogenous variables. *International Journal of Forecasting* 33(3), 627–651.
- [20] Papana, A., C. Kyrtsou, D. Kugiumtzis, and C. Diks (2017). Financial networks based on granger causality: A case study. *Physica A: Statistical Mechanics and its Applications* 482, 65–73.
- [21] Rogers, B. P., V. L. Morgan, A. T. Newton, and J. C. Gore (2007). Assessing functional connectivity in the human brain by fmri. *Magnetic resonance imaging* 25(10), 1347–1357.
- [22] Song, H. and Z. Zhang (2014). Analyzing multiple multivariate time series data using multilevel dynamic factor models. *Multivariate Behavioral Research* 49(1), 67–77.
- [23] Uematsu, Y. and T. Yamagata (2025). Discovering the network granger causality in large vector autoregressive models. *Journal of the American Statistical Association* 120(552), 2385–2396.
- [24] van de Geer, S., P. Bühlmann, Y. Ritov, and R. Dezeure (2014). On asymptotically optimal confidence regions and tests for high-dimensional models. *The Annals of Statistics* 42(3), 1166 – 1202.
- [25] Wainwright, M. J. (2019). *High-dimensional statistics: A non-asymptotic viewpoint*, Volume 48. Cambridge University Press.
- [26] Wei, W. W. (2019). *Multivariate time series analysis and applications*. John Wiley & Sons.
- [27] Zhang, C.-H. and S. S. Zhang (2014). Confidence intervals for low dimensional parameters in high dimensional linear models. *Journal of the Royal Statistical Society Series B: Statistical Methodology* 76(1), 217–242.
- [28] Zhang, J. and Y. Li (2023). High-dimensional gaussian graphical regression models with covariates. *Journal of the American Statistical Association* 118(543), 2088–2100.

- [29] Zhu, K. and H. Liu (2022). Confidence intervals for parameters in high-dimensional sparse vector autoregression. *Computational Statistics & Data Analysis* 168, 107383.
- [30] Zou, H. (2006). The adaptive lasso and its oracle properties. *Journal of the American statistical association* 101(476), 1418–1429.